

Markov Chain Attribution + Budget Optimization — Deployment Design (Brief)

1) Purpose & Scope

Deploy a weekly batch pipeline that:

- Builds **Order-1 Markov** attribution from customer journeys.
- **Calibrates** a blended weight $\text{opt_weight} = \alpha \cdot \text{removal_weight} + (1-\alpha) \cdot \text{provided_weight}$ (α via SMAPE minimization).
- Solves a **linear program** with $\sum x == \text{budget}$ under business/strategic constraints.
- Publishes **recommendations & artifacts** for BI and stakeholders.

Primary KPI: predicted conversions at recommended mix

Secondary KPIs: uplift vs. current, attribution stability (MoM), runtime < 10 min.

2) Inputs & Outputs

Inputs (read-only, PostgreSQL `ads_db5.dw4`)

- `customer_journeys(customer_id, journey_step, channel, timestamp, converted, conversion_value, customer_segment, ...)`
- `channel_performance(channel, current_monthly_spend, conversions_attributed, cost_per_conversion, min_monthly_spend, max_monthly_spend, conversion_rate_per_1000)`

Artifacts (written to storage; optional DB publish)

Path / Table	Description
<code>processed/transition_matrix.csv</code>	Absorbing Markov transition matrix (Order-1)
<code>processed/markov_attribution.parquet</code>	Removal effects + normalized weights (negatives floored to

<code>processed/channel_performance_enriched.parquet</code>	Performance + <code>opt_weight</code> (α -calibrated)
<code>processed/budget_scenario_{BUDGET}.parquet</code>	Optimal spend per channel for a given budget
<code>processed/sweep_alpha_budget.parquet</code>	Grid of predicted conversions across $\alpha \times$ budget
(Optional) <code>analytics.budget_scenarios</code>	DB table for dashboards (scenario_id, budget, channel, optimal_spend, predicted_conversions, spend_change_pct)

3) Processing Logic

[Postgres ads_db5.dw4: customer_journeys (RO), channel_performance (RO)] \rightarrow (SQLAlchemy/pandas) \rightarrow [ETL & Attribution: normalize channels; build transition matrix (Order-1); baseline conversion via fundamental matrix; removal effects (floor negatives; normalize); α calibration (SMAPE vs historical)] \rightarrow [Optimization (LP, CBC): objective = maximize $\sum (x/1000) \cdot conv_rate \cdot opt_weight$; constraints = $\sum x == BUDGET$; $min \leq x \leq max$ per channel; strategic (Paid $\geq 30\%$ budget; Social+Display $\leq 40\%$); outputs = optimal_spend, predicted_conversions, binding flags] \rightarrow [Publish Artifacts: Parquet/CSV in ./processed/; optional upsert to `analytics.budget_scenarios`] \rightarrow [Notify: Slack/Email summary (α^* , uplift, bindings, runtime)]

4) Orchestration

- **Schedule:** Weekly (e.g., Monday 06:00 ET) after data closes.
- **Trigger:** Cron / Airflow / GitHub Actions.

Cron example (server time ET):

5) Implementation Notes

Runtime stack - Python 3.11; `pandas`, `numpy`, `sqlalchemy`, `psycopg2-binary`, `python-dotenv`, `pulp` (CBC). - Containerized (Docker) for reproducibility.

Secrets & Config - `.env` or Vault: `POSTGRES_USER`, `POSTGRES_PASSWORD`, `POSTGRES_HOST`, `POSTGRES_PORT`, `POSTGRES_DATABASE`. - Config: `budgets`, `alpha_mode={"calibrated"|"fixed"}`, `alpha_fixed`, constraints toggles, notify webhook.

Processing steps 1. **Extract**: set `search_path=dw4`; read both tables to DataFrames (read-only role). 2. **Transform & Attribution**: - Title-case/normalize channels; build absorption matrix; compute baseline conversion. - Removal effect per channel; floor negatives to 0; normalize to weights. 3. **Calibration**: grid $\alpha \in [0,1]$; minimize SMAPE between predicted @ current spend and `conversions_attributed`; set `alpha_best`. 4. **Optimization**: for each budget, solve LP with $\sum x == \text{budget}$ and business bounds; tag binding constraints; compute predicted conversions. 5. **Publish**: write artifacts to `./processed/`; optional DB upsert to `analytics.budget_scenarios`. 6. **Notify**: summarize α^* , uplift vs current, binding constraints, elapsed time.

6) Quality, Monitoring, Alerts

Data checks - Required columns present; `journey_step ≥ 1`. - Transition matrix row sums ~ 1.0 ($\pm 1e-6$). - Baseline conversion $\in [0,1]$. - Negative removal effects captured & floored; list logged.

Calibration - Report α^* and SMAPE; alert if SMAPE > 0.35 .

Optimization - Solver status = Optimal; $\sum x == \text{budget}$. - Report **binding** min/max/strategic constraints.

Drift (optional) - MoM drift of weights per channel; alert if $> 50\%$ relative change.

Alerts - Job failure, SLA breach, calibration drift threshold, solver infeasible.

7) Security & Compliance

- Principle of least privilege: RO to `dw4`; separate RW user only for `analytics` (if used).
 - Secrets via env/Vault; no secrets in logs; artifacts contain no PII.
 - Access logs retained per org policy.
-

8) Rollback & Reproducibility

- Each run writes `processed/YYYYMMDD_HHMMSS/` .
 - Store: git SHA, α^* , budgets, constraints, solver log.
 - Rollback by pointing BI/API to last good timestamp or DB snapshot.
-

9) Interfaces

BI/Dashboard - Read Parquet/CSV or `analytics.budget_scenarios` . -
Standard visuals: Current vs Optimal spend, Predicted conversions,
Binding constraints, Sensitivity ($\alpha \times$ Budget).

(Optional) Read-only API - `GET /recommendation?budget=65000` →
JSON: per-channel optimal_spend, predicted_conversions, α^* , uplift,
binding flags.